



SPECIAL TOPIC ARTICLE

Trustworthy residual vehicle value prediction for auto finance

Mihye Kim¹ | Jimyung Choi¹ | Jaehyun Kim¹ | Wooyoung Kim¹ |
Yeonung Baek¹ | Gisuk Bang¹ | Kwangwoon Son¹ | Yeonman Ryou¹ |
Kee-Eung Kim²

¹Hyundai Capital Services, Seoul, South Korea

²Kim Jaechul Graduate School of AI, KAIST, Daejeon, South Korea

Correspondence

Mihye Kim, Hyundai Capital Services, Korea.

Email: mee9110@gmail.com

Abstract

The residual value (RV) of a vehicle refers to its estimated worth at some point in the future. It is a core component in every auto finance product, used to determine the credit lines and the leasing rates. As such, an accurate prediction of RV is critical for the auto finance industry, since it can pose a risk of revenue loss by over-prediction or make the financial product incompetent through under-prediction. Although there are a number of prior studies on training machine learning models on a large amount of used car sales data, we had to cope with real-world operational requirements such as compliance with regulations (i.e., monotonicity of output with respect to a subset of features) and generalization to unseen input (i.e., new and rare car models). In this paper, we describe how we addressed these practical challenges and created value for our business at Hyundai Capital Services, the top auto financial service provider in Korea.

INTRODUCTION

In auto finance, a vehicle is the main object or purpose for which the financial service is provided. As such, the estimation of the residual value (RV) of the vehicle is a critical factor in determining the credit lines and the leasing rates. As an example, for an auto loan, the estimated market price of the secured vehicle is used to determine the credit line. As another example, for an auto lease, the predicted future value of the vehicle at the time of its return is used to determine the monthly rate. If the RV of the leased vehicle is over-predicted than the actual value at the time of return, it would impose a risk of revenue loss, while an under-prediction will push up the lease rate and thereby potentially lose the customer to competitors. Thus, an accurate prediction of RV is one of the most

important business requirements of every auto financial service.

Accurate RV prediction is also critical for the used car market as a whole, beyond the auto finance business. In most countries including Korea, the used car market is commonly recognized as a *lemon market*, where information asymmetry exists between sellers and buyers. Buyers lack sufficient information for determining a fair price for the vehicle and thus worry about paying more than what it is worth. This leads to buyers making their best attempt to pay less, and auto dealers consequently sell cars of inferior quality to match the price. Ultimately, both the sellers and the buyers become dissatisfied with the used car market, resulting in a decreasing number of market participants. This unfortunate market behavior is difficult to rectify unless the information balance between sellers and buyers

This is an open access article under the terms of the [Creative Commons Attribution-NonCommercial-NoDerivs](https://creativecommons.org/licenses/by-nc-nd/4.0/) License, which permits use and distribution in any medium, provided the original work is properly cited, the use is non-commercial and no modifications or adaptations are made.

© 2023 Hyundai Capital Services and The Authors. *AI Magazine* published by John Wiley & Sons Ltd on behalf of Association for the Advancement of Artificial Intelligence.

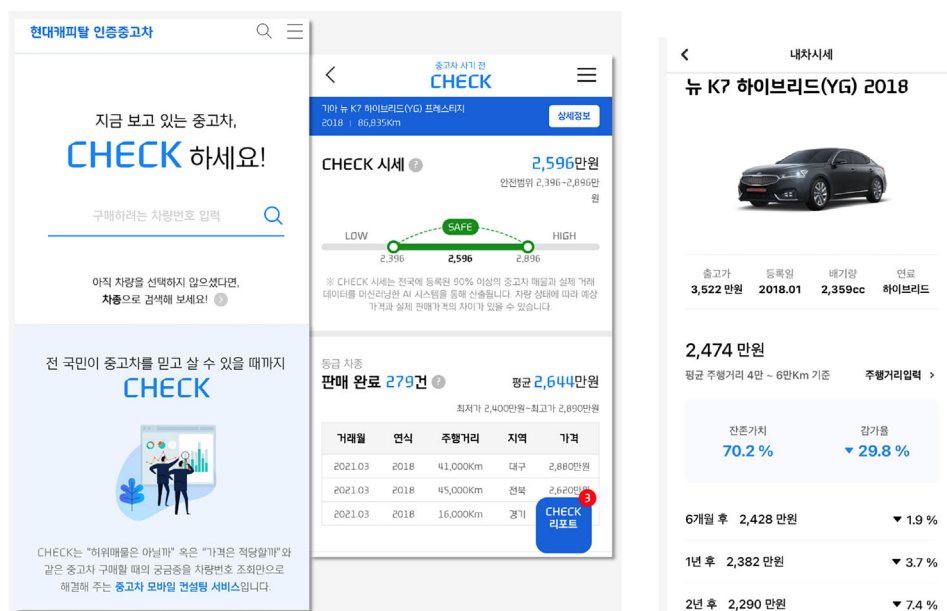


FIGURE 1 Screenshots of the used car price estimation service menu in the HCS mobile app, for potential buyers (left) and potential sellers (right). The price estimation can be made at the individual car level by supplying the license plate number.

is restored. This is also one of the goals behind our work, to make accurate predictions of used car prices available to ordinary consumers to protect them against exploitative dealers, thereby contributing to a healthier and more robust market for used cars. This ultimate goal is also important for us at Hyundai Capital Services (HCS), the top auto financial service provider in Korea ever since its launch of the retail (installment) finance business in 1996.

As a part of our effort, we have been providing a personalized used car price estimation service through a mobile channel since 2018, shown in Figure 1, to improve customer engagement. This service leverages a deep neural network (DNN) model trained on proprietary used car sales data. At the time of the service launch, it was the first to adopt deep learning and thus attracted a lot of attention. Yet, the model was unsuited for the core business, since it fell short in many operational requirements: as shown in the left pane of Figure 7, although the model captures the general trend of value depreciation with respect to mileage and age of a specific car model, the monotonicity is clearly violated in many local regions, as can be seen from the spikes and valleys. In addition, the model was not able to handle novel or rare car models. This paper describes how we addressed these practical limitations in the baseline model and created business values by integrating the improved model into the core component of auto finance services, handling tens of thousands of monthly cases and increasing quarterly profit by several million US dollars since Q3 2022.

RELATED WORK

The vehicle RV prediction can be naturally formulated as a regression task in supervised learning, using many sources of used car sales transaction data readily available. There is thus a vast amount of literature on the topic, ranging from classical linear regression models to deep neural networks, and we only highlight some of the most relevant work for our approach.

Most notably, Lessmann and Voß (2017) report a systematic study on empirical performances of 19 regression models on a proprietary dataset consisting of 4.5 million instances covering 6 car models from the same manufacturer. The regression models used for experiments range from classical linear regression models, non-linear models, and ensemble models. The study finds that non-linear models and ensemble models perform much better than linear models, which is a natural result given that the value depreciation is inherently non-linear. This finding is consistent with other studies on RV prediction, for example, support vector machine regression (Lessmann, Listiani, and Voß 2010) and neuro-fuzzy systems (Lian, Zhao, and Cheng 2003; Wu, Hsu, and Chen 2009). More recent work adopts an advanced training paradigm to learn a good representation of input for improving RV prediction accuracy (Rashed et al. 2019).

However, non-linear models are often unexplainable. In highly regulated domains such as finance, the model needs to be fully audited before its deployment. Due to

this operational requirement, linear regression models are still popular in practice. The AutoCycle by Moody's analytics uses logistic regression and quantile regression models to provide a commercial service for vehicle value prediction (Hughes et al. 2016). It allows forecasting of car prices under various macroeconomic scenarios such as changes in the unemployment rate and gas prices. The model was reported to achieve an R-squared of 0.89 on 6.4 million test instances. Yet, being a linear prediction model by its nature, it relied on heavily engineered features that require substantial expert knowledge, despite a huge amount of data (approximately 31 million instances from the National Automobile Dealers Association) being available for training and validation.

Our work focuses on training an expressive non-linear model for RV prediction from a large amount of data with the following desiderata: (1) enforce operational requirements of the model by introducing monotonicity constraints (e.g., RV prediction should be monotonically decreasing with mileage and age) for fairness and transparency, and (2) make the model generalize well on out-of-distribution data (e.g., new car models released to the market by the manufacturers). To the best of our knowledge, these two requirements are yet to be simultaneously addressed for a commercial-grade machine learning model, at least for the RV prediction task of vehicles in auto finance.

TASK DESCRIPTION

Our objective was to develop an expressive (i.e., non-linear) machine learning model that accurately estimates vehicles' current and future market prices while addressing the following operational requirements:

1. For the same car model, the RV ratio (i.e., the price-to-Manufacturer Suggested Retail Price(MSRP) ratio) should be lower if the mileage (total distance traveled) is higher and if the age is older. Also, the RV ratio should be higher for more recently released models and for lower-level trims. The financial authorities regulate whether the finance companies' loan limits or pricing standards are adequate and well-explained to protect financial consumers from undue damages. As a vehicle price prediction model establishes standards for setting the financial product's credit lines and pricing in the auto finance industry, the model used needs to be highly accurate as well as have a clear rationale for its logic.
2. The model should be able to estimate RVs for all car models including those new on the market with no history of used car trading. The actual used car sales data is the main source of information for training the model, but new car models are continuously released quarterly

or semi-annually, and these newer models inevitably lack sufficient sales data for reliable training. Financial companies, nonetheless, face the challenge of having to estimate the market price of fresh models even before there is any data available to add to the training set, since the lease product requires predicted RV at the end of the lease to determine monthly payments.

3. The model needs to be updated continuously and regularly. Since the economic situation is always changing, more recent used car sales data is more accurate reflection of how much the car is worth today, compared to older ones. It is important to maintain the prediction accuracy by regularly obtaining the latest data and updating the model.

The model is trained to predict the RV ratio, that is, the ratio of the used car price to the MSRP of the new car. As for the training dataset, we secured about 1.8 million used car sales transaction records from 600 vehicle models which were actually traded in Korea over the past 12 years. Each transaction record contained a number of attributes such as vehicle specifications (manufacturer, brand, model, new car price, etc.), vehicle usage information (mileage, age, whether it was for rental, etc.), and market supply & demand (no. of used cars available on the market for sale, new car sales volume, used car searches on platform, time taken for sale, etc.). There were a total of 71 candidate features, including 40 vehicle specifications, 26 vehicle usage information, and 5 market condition information. Additional information on the summary statistics of the training data is provided in Figure 2.

METHODOLOGY

Dataset

We used a proprietary dataset of used car sales transactions, published monthly by the dealers' association in Korea. We used the actual sale prices as the target of the model for training. We filtered out instances that were obviously incorrect or beyond the scope of our interest. For example, transactions of commercial vehicles such as taxis and trucks, RV ratio greater than 1.0, and yearly mileage greater than 200k km (124k miles) were removed, which comprised 6% of the original data.

Anomaly detection

Removing anomalies is a crucial process in exploratory data analysis (EDA) for improving a model's prediction accuracy. It was even more important in our case because most vehicle sales price data were entered manually by




FIGURE 2 Statistics chart of training dataset.

dealers. This introduces opportunities for human errors and even fraudulent cases such as forged sales price to avoid taxes. Thus it was essential to filter out abnormal instances in the dataset before training the RV prediction model. Although there are many promising anomaly detection methods, we limited ourselves to comparing Extended Isolation Forest (EIF) and Local Outlier Factor (LOF) which do not rely on any distributional assumption.

The EIF method is a tree-based algorithm effective for multi-dimensional data with large volumes. It calculates the leaf node (terminal node) distance for each instance to become isolated to derive the score. The score is higher if the distance (depth) is shorter, and the outliers are defined based on the score value (Hariri, Kind, and Brunner 2019). The LOF method defines anomaly points by taking into consideration the relative density among the clusters. It calculates the local density based on k-nearest neighbors (KNN), and compares the local density of one object against the local density values of its neighbors, categorizing those with large gaps as outliers (Breunig et al. 2000).

We measured the precision and recall on a subset of data, restricted to one particular vehicle model with the largest sales volume in the new and used car market. For this subset of data, we recruited domain experts to label the instances as normal versus abnormal to compute the performance metrics. This held-out labeled dataset contained approximately 60k instances. When comparing the performances of the two methods at the same threshold (0.3%), EIF performed better than LOF (Figure 3).

In addition, the outlier detection performance of the two algorithms is compared to a more traditional method, IQR, which is a method of setting the range for each major feature (age, mileage) using the interquartile range, and then defining the object outside the range as an outlier (Vinutha, Poornima, and Sagar 2018). EIF outperformed IQR in anomaly detection in individual vehicle models while LOF showed a tendency to classify some of the normal data as outliers (Figure 4). Therefore, we adopted EIF for carrying out multi-dimensional anomaly detection.




FIGURE 3 Performance comparison of anomaly detection algorithms.

Vehicle model code vectorization leveraging DNN

There are about 500 vehicle models that need to be covered by our RV prediction model, some of which have sufficient data while others do not. It was one of the important requirements for us to predict the RV as accurately as possible for all vehicle models, including the ones that lacked sufficient amount of training data. For this purpose, it was required to vectorize the “Vehicle Model Code,” a categorical variable that distinguishes individual vehicle models, and use it on the prediction model. Even though the generally used one-hot encoding converts the categorical values to numerical values, it had limitations in reflecting the unique characteristics of each vehicle model.

Therefore, we employed the embedding method. Embedding has the advantages of generating denser




FIGURE 4 Comparison of anomaly detection algorithms. (Black Dots: Outliers classified by each method.)




FIGURE 5 Neural network architecture for obtaining vehicle model embedding.

representations compared to one-hot encoding, and showing not only characteristics of each vehicle model but also similarities between vehicle models with finite-dimensional vectors.

The embedding was obtained by plugging an embedding layer into a neural network, trained to predict RV ratio (Figure 5). Embedding dimension was set to 2 for the sake of interpretability and visualization for domain experts. We considered various architectures for the neural network (e.g. number of layers and hidden units), and the first batch of candidates was selected solely based on the RV ratio prediction accuracy of the neural network. They were then reviewed by domain experts, and further narrowed down by visually inspecting whether the overall patterns of embedding was consistent with their knowledge. The final embedding was selected by the RV ratio prediction accuracy of the CatBoost model.

Using only two dimensions for embedding, the performance of the RV prediction model was significantly improved. When the embedding vectors of vehicles were visualized, they could be clustered in a meaningful way (Figure 6). We used this clustering result to assign embed-

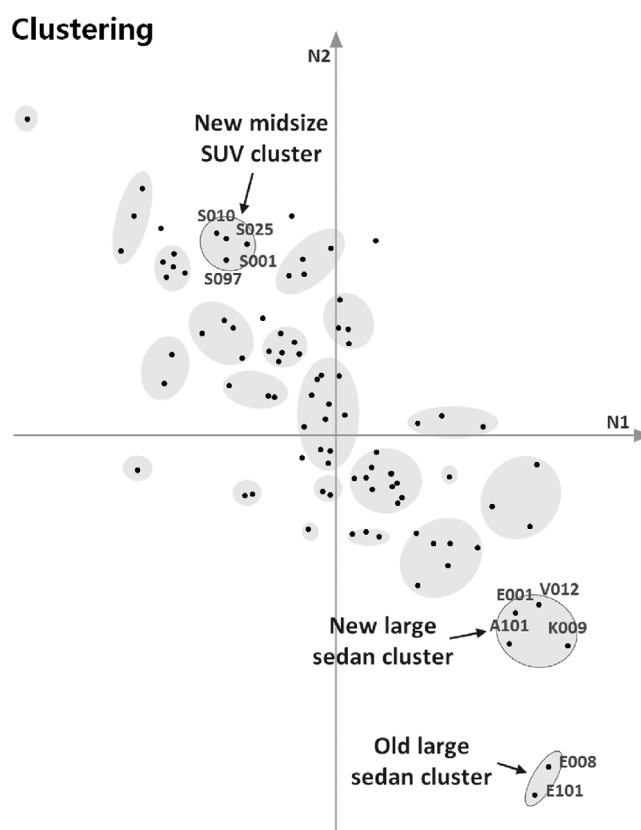


FIGURE 6 Visualization of vehicle embedding vectors. Vehicle models are anonymized.

ding vectors to vehicles that could not be projected onto embedding vector space due to a lack of data. For example, if the specification information of a newly released vehicle is similar to that of vehicles in a specific cluster, the centroid of the corresponding cluster is assigned as an embedded vector value of the new vehicle. For a car model with complex characteristics that cannot be determined by a single cluster, we used two or more selected clusters. Using the centroids of each cluster, we found a new point that reflects the cluster information and use it as an embedding vector. By doing so, we

were able to solve the cold-start problem caused by the lack of data on newly launched and rarely traded car models.

Training boosted tree models

We selected tree-based models for the following reasons: First, we aimed to improve the prediction accuracy for vehicle data by capturing complex complicated non-linear relationships between variables. In addition, the model should be able to conveniently reflect the monotonicity of the target variable with respect to a subset of features. We trained the models using XGBoost, one of the most powerful and popular tree-based machine learning methods, and CatBoost, which specializes in learning categorical variables. In particular, CatBoost has the advantage of mitigating data leakage and recognizing categories with similar characteristics in groups because it processes categorical features via the mean target encoding method (Dorogush, Ershov, and Gulin 2017). This is particularly helpful for the vehicle data since they contain various categorical features such as vehicle model code, color, and maker. We compare the performances of XGBoost and CatBoost in this section.

We trained the model by setting the RV ratio as the target variable. From a total of 71 features, we determined that 38 of them being of low relevance or replaceable by other derivative features. Another 20 of them had too many missing value cases. We thus selected 13 features from the dataset, and augmented them with 8 additional features from side information (e.g., how long since the car model generation was first released in the market), resulting in a total of 21 features. The selected features were 5 vehicle specifications, 11 vehicle usage information, and 5 market demand/supply information.

As mentioned in the Task Description section, we set the constraints checking the monotonicity when creating a tree, to satisfy the monotonic relationship between a subset of features and the RV ratio. Monotonic constraints were given to four features; age (month), mileage, new car price, and model year. In the model that did not account for monotonicity, the price prediction shows significant fluctuation depending on the age and mileage, as can be seen in the left graph of Figure 7. On the other hand, in the model we developed with monotonicity constraint, price prediction shows monotonicity for specific variables, as displayed in the graph on the right. Monotonicity enforces the model to accord with common sense; the more high-end options the vehicle has, and the more aged the vehicle is, the faster its depreciation. Also, the encoded vectors of vehicle model code derived by leveraging the DNN model were used as input features for

training the boosted tree models. The final model consists of approximately 4500 trees, with the maximum depth of 9.

RESULTS AND DISCUSSION

We conducted a performance evaluation of the vehicle price model for out-of-time (OOT) data sets, and used the mean absolute error (MAE) of the actual RV ratio and model-predicted RV ratio as evaluation metrics. As we know, MAE is defined as $\sum_{i=1}^n |y_i - x_i|/n$ where y_i is the predicted value, x_i is the observed value and n is the total number of data points. Also, by defining $MAE_{car-seg}$ as below, the offset effect for each car model was aggregated in calculating MAE:

$$MAE_{car-seg} = \frac{\sum_i |\overline{y_{car_i}} - \overline{x_{car_i}}|}{n_{car}} \quad (1)$$

where $\overline{y_{car_i}}$ is the average predicted value of car model i , $\overline{x_{car_i}}$ is a mean observed value of car model i , and n_{car} is the total number of car models in the data set. The performance of models that were developed using each algorithm is shown in Table 1.

Compared to the baseline model (DNN only), currently deployed for used car estimation service in Figure 1, the new models showed a substantial improvement in error rate while satisfying the monotonicity for key variables complying with regulations; a 44% error reduction compared to the baseline model by XGBoost, and a 47% error reduction by CatBoost. There were even more improvements in the error rate by car segment than the total errors of individual vehicles (a 76% error reduction by XGBoost, an 80% error reduction by CatBoost). It is due to the fact that the characteristics of the vehicle model were reflected in the models through embedding vectors. It was also confirmed that CatBoost outperformed XGBoost (Table 2). This is a natural result, given that CatBoost handles categorical features better than XGBoost, and our training data contained many categorical features. Therefore, we selected DNN+CatBoost as the final model. In the final model, the overall errors decreased 47% and errors by vehicle model decreased 80% compared to the baseline model (DNN only) while satisfying the monotonicity of key features (Figure 7).

Lastly, to verify the effectiveness of removing anomalies, we used raw data to train the prediction model leveraging the DNN+CatBoost model. When compared to the model trained with raw data, the model trained with cleansed data preprocessed using anomaly detection decreased MAE by 0.87%, from 4.51% using the raw data to 3.64% using the cleansed data.




FIGURE 7 Comparison of predicted RV ratios projected on age and mileage per year.

TABLE 1 Comparison of vehicle price prediction accuracy between the baseline DNN model, and the boosted tree models (XGBoost and CatBoost).

Model	MAE (%)		Coverage	Monotonicity
	Total	Car-Seg		
Baseline DNN	6.87	6.55	77%	Unsatisfied
DNN + XGBoost	3.82	1.62	100%	Satisfied
DNN + CatBoost	3.64	1.32	100%	Satisfied

Abbreviations: DNN, deep neural network; MAE, mean absolute error.

TABLE 2 Performance by vehicle model of the prediction models leveraging XGBoost (XGB) and CatBoost (CAT).

Car segments	MAE(%)	
	DNN+XGB	DNN+CAT
SUV/RV	3.24	3.15
Large-size	3.01	2.92
Mid-size	3.65	3.55
Compact	4.17	4.03
Mini-compact	4.24	4.01
Mini-truck(commercial)	4.91	4.67
Minivan	5.96	4.98

Abbreviations: DNN, deep neural network; MAE, mean absolute error.

In summary, we were able to develop a new model with significantly improved accuracy. The baseline model could not be applied for our company's core business processes (calculating the lien value for used car loans, predicting vehicle return prices to set the monthly lease rates, and leveraging reserves) due to its issues in the prediction accuracy, the cold-start problem with new car models, and the

lack of monotonicity guarantee for key features. It was therefore only used for the casual used car price estimation service available through our company's mobile channel. The new model resolves all three issues above, improving prediction accuracy by at least 15% for all car models. The business unit thus approved our model for deployment in early 2022.

APPLICATION USE AND PAYOFF

First, we used the model to determine the residual value of our leased vehicles. As mentioned at the beginning of this paper, an incorrect estimate of residual value can have a direct impact on our profit and loss. At the end of a lease, customers can choose to return their car or buy it out. If the residual value we have estimated is higher than the market value, it is more likely for the customer to return the car. When we end up selling the returned car, the difference between the predicted value and actual market value would be added to our loss. To reduce this loss, we could

simply set the residual value at the end of the lease as low as possible, but this would have made monthly lease payments expensive, resulting in less consumers deciding to lease a car with our company. The more accurately we can predict the residual value of a returning car, the less the loss caused by a prediction error. By replacing the baseline model with the one that we developed, we were able to offset our loss caused by an error in the predicted RV by approximately 21%. The gap between the forecast value and the actual price at sale was reduced to half the original level. This enabled us to reduce our reserves by a significant amount (several millions of dollars in Q3 2022 alone).

Second, our model was used as the benchmark in determining the credit line on loans provided for buying a used car or those secured by a used car. By offering a credit line that approximates the value of the car, we can reduce the account's probability of default, and also prevent fraud cases where someone may take out a loan for an amount greater than what the vehicle is actually worth and intentionally default on the loan. Another benefit is that that we can encourage customers to spend reasonably, within their ability to repay.

Third, we used the model to offer a service allowing customers to find out what their car is worth on our existing mobile application, "Hyundai Capital." Car prices estimated from the previous, DNN-based model were replaced with the ones from the newly developed model. This enabled us to offer more accurate forecast values to mobile application users. By using this service when planning to sell their car, customers can protect themselves from being taken advantage of by shady dealers. The service also helps customers decide to return the car at the end of a lease. We expect our enhanced service will attract more customers to our app and build stronger customer trust.

Lastly, we are selling the vehicle price database generated from our RV prediction model to used car online market companies via partnership. We have entered into contracts where we provide the partner companies with prices predicted by our model in return for lowering our advertising costs by 10-20%. We are steadily expanding these partner companies as we further solidify our position as the market leader in price forecasting.

Other than the areas of our business, we also use the model to better protect consumers and improve transparency in the used car market. As mentioned in the introduction, the problem of information asymmetry in the used car market has often made used car buyers susceptible to unfair deals. To prevent this, the South Korean government has established a website hosting a standard platform for used car companies (www.car365.go.kr). Used car prices estimated by our company are one of the standard prices shared on the government website. Aside from sharing our prices through the government agency, we are

also engaged in our own efforts to further assist consumers. A case in point is our service named "What to check before buying a used car" available on our Certified Pre-Owned car site. As its name suggests, the service helps consumers check before the purchase whether the car they have their eye on does actually exist, and whether the car is worth the price tag. We are enhancing the trustworthiness of our car price forecasts through the improvement we have made to the model, and thereby contribute to stronger protection of consumer rights.

In addition, to maintain a robust and healthy market for used cars, the Korean government requires validation when finance companies issue used car loans. The car prices predicted by our model have provided resources for setting the standard in this regard. With its enhanced accuracy, the newly launched RV model is expected to aid credit line-related risk management in the entire used car market.

DEPLOYMENT AND MAINTENANCE

In the first quarter of 2022, we developed the model. In the second quarter, we deployed it into the HCS internal system and reviewed which business areas required RV forecasting. Finally, in the third quarter, we applied the model to core business processes. In the fourth quarter of 2022, we created a dashboard that allows us to monitor the database values generated by the model at a glance. This has enabled robust monthly operation of the model. Currently, the new model is used to handle several tens of thousands of customers each month for lease and used car loan products. Additionally, a similar number of customers are using the 'Check out My Car's Value' service on Hyundai Capital's mobile channel.

To maintain the performance and accuracy of our RV prediction model, we have designed a workflow to retain it with the latest transaction data every month. We have built the model's operations architecture by setting up a vehicle data mart on our internal big-data platform and developing automated process modules (Figure 8). Vehicle data marts are composed mainly of the transaction data from the used car market and sales data of vehicles whose lease contracts have matured. Market information such as oil prices and vehicle shipments is added as well. These are used for retraining the model. The model development process described above was also modularized so that the model was automatically retrained in the modeling server. The model produces two outputs, the used cars' current price estimates and their future residual value predictions. The current price data is updated monthly on the Hyundai Capital price inquiry service platform through a relay server. Future price data is loaded through

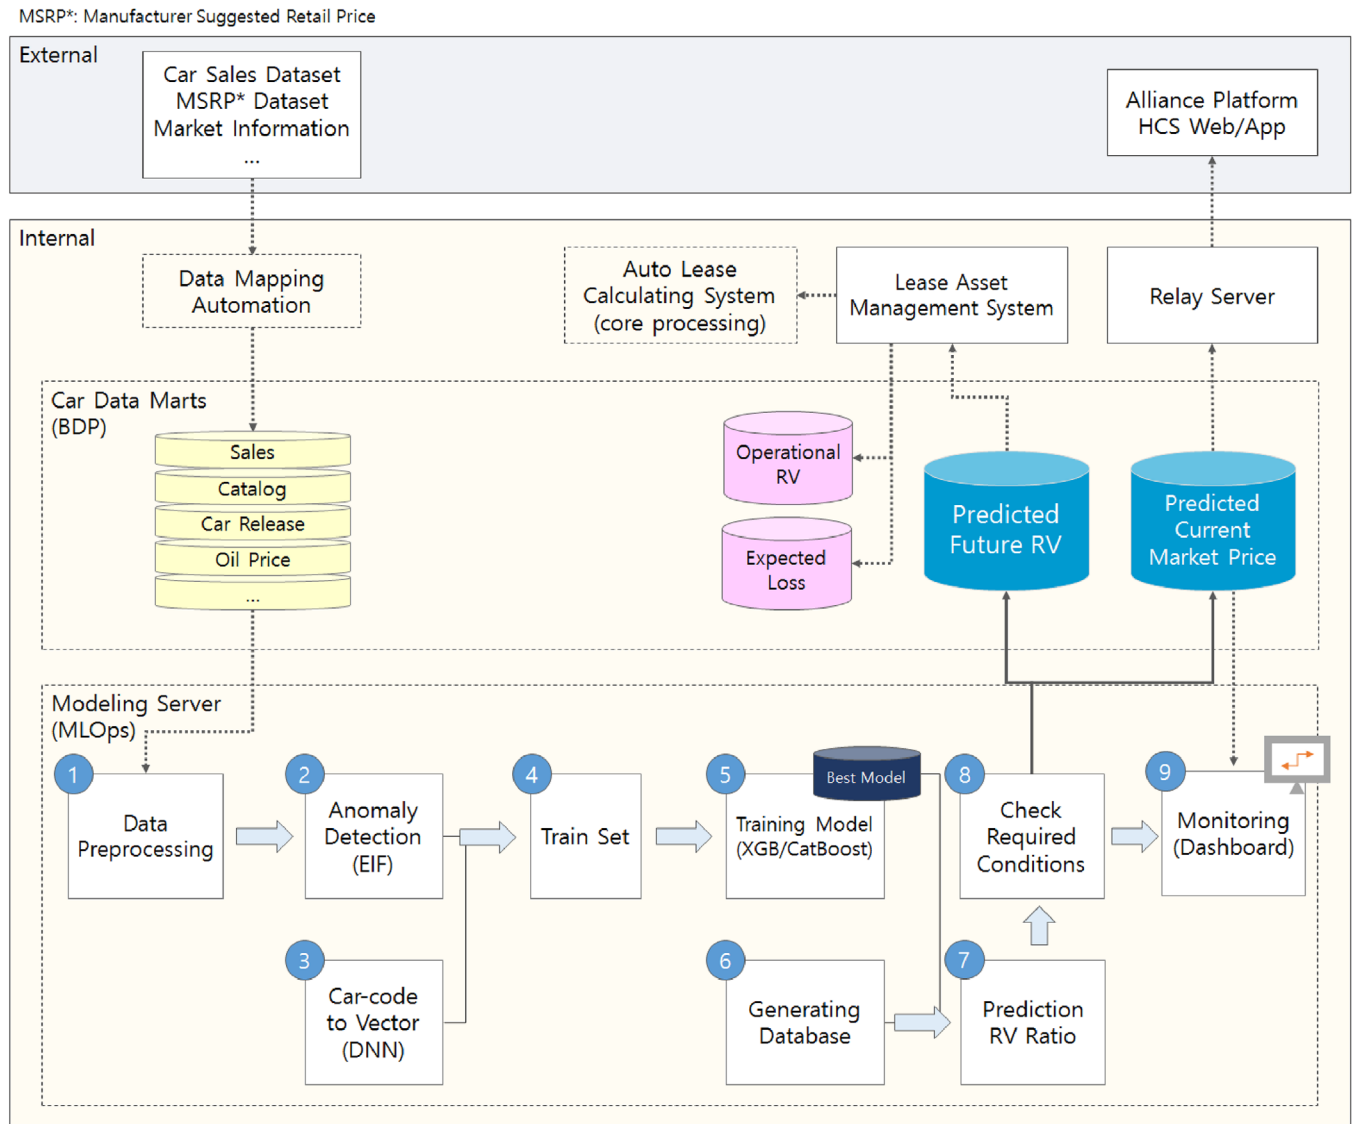


FIGURE 8 A workflow diagram illustrating automated operations of our vehicle price predicting model.

the internal strategic management system, and used for offering automotive lease products to customers.

Once the RV prediction model is retrained, it goes through a safety measure that automatically monitors the calculated RV ratios to review whether they are appropriate before delivering them to the relevant systems. To make the monitoring results more accessible, a dashboard was designed in the part of MLOps system (Figure 9). Through this, not only the risk modeling team but also the risk strategy team, marketing team, and sales department can immediately check the model validation results. It was confirmed that the model is in a suitable state to operate without any issues, based on the several month of monitoring result (Table 3).

The quality control check process is as follows. There are two types of check points: model performance validation and operational stability check (health check). Model val-

TABLE 3 The performance of the model over the last three months. The baseline performance of the model was calculated in the first quarter of 2022.

	Baseline	Feb.2023	Mar.2023	Apr.2023
MAE(%)	3.64	3.73	3.72	3.76

Abbreviation: MAE, mean absolute error.

idation uses MAE as an evaluation index to check model performance and its level of variation from the previous month. In addition, the monotonicity is checked based on essential variables such as age, mileage, and trim that ensure the monotonicity of the residual price within the same vehicle model. The stability check confirms the level of change in the RV ratio by vehicle type compared to the previous month, as well as the appropriateness of the vehicle depreciation. The change in the number of

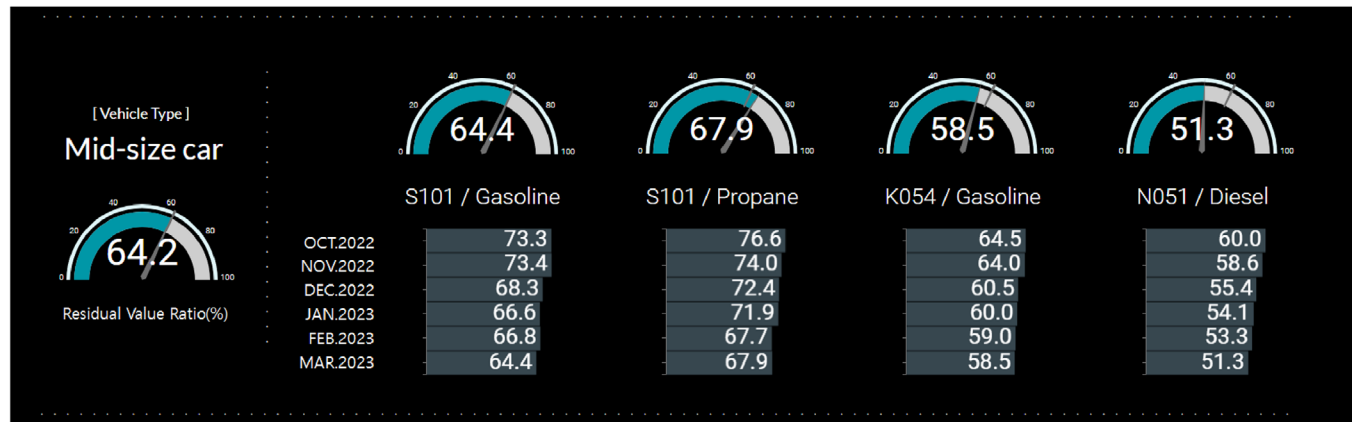


FIGURE 9 This is an example of the dashboard used for monitoring under the model validation step. The example shows the average residual value for the mid-size car group and individual vehicle models within the group.

vehicles whose market price can be calculated is also checked against the total number of vehicle types traded in the market.

CONCLUSION

Predicting vehicle RV is a critical factor in the auto finance business, and for this reason we dedicated our efforts on enhancing the accuracy of the prediction model. While the single most important goal for our research was to improve the prediction performance of our model, to be able to apply it to a wider range of business areas, we also aimed to satisfy the monotonicity of key features as well as overcome the issue of data shortage for some models (cold-start problem). To address these challenges, we sought to upgrade model development across its entire stages, and enhanced the model by applying new techniques. We leveraged anomaly detection to sort out the labels, and also applied neural network to ensure embedding of the vehicle model. We also trained tree-based machine learning methods to achieve a significant jump in the model's prediction accuracy (overall errors decreased 47%, errors by vehicle model decreased by 80%), at the same time as resolving the monotonicity and cold-start problems.

Through this research, we have increased the reliability of our new model and successfully applied it to our core business, resulting in meaningful business outcomes. As the leading auto finance service provider in Korea, we are dedicated to providing reliable used car prices to the market and protecting the rights of our customers. This research has enabled us to further contribute to this mission. Moving forward, the team at Hyundai Capital Service will continue to explore and leverage the latest AI techniques to maintain and improve the accuracy of our residual value predictions. Specifically, we are currently

developing a time-series prediction model that focuses on macroeconomic and market transaction information to forecast short-term fluctuations in vehicle market prices. In addition, we plan to conduct research on personalized residual value prediction models that take into account the characteristics of individual drivers.

CONFLICT OF INTEREST STATEMENT

The authors declare that there is no conflict.

ORCID

Mihye Kim <https://orcid.org/0009-0008-9354-9364>

REFERENCES

- Breunig, M. M., H.-P. Kriegel, R. T. Ng, and J. Sander. 2000. "LOF: Identifying Density-Based Local Outliers." In *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data*, 93–104.
- Dorogush, A. V., V. Ershov, and A. Gulin. 2017. "CatBoost: Gradient Boosting with Categorical Features Support." In *Workshop on ML Systems at NIPS*.
- Hariri, S., M. C. Kind, and R. J. Brunner. 2019. "Extended Isolation Forest." *IEEE Transactions on Knowledge and Data Engineering* 33(4): 1479–89.
- Hughes, T., S. W. Malone, M. Brisson, and M. Vogan. 2016. "Introducing AutoCycleTM: Residual Risk Management and Lease Pricing at the VIN Level." Technical report, Moody's Analytics.
- Lessmann, S., M. Listiani, and S. Voß. 2010. "Decision Support in Car Leasing: A Forecasting Model for Residual Value Estimation." In *International Conference on Information Systems*.
- Lessmann, S., and S. Voß. 2017. "Car Resale Price Forecasting: The Impact of Regression Method, Private Information, and Heterogeneity on Forecast Accuracy." *International Journal of Forecasting* 33(4): 864–77.
- Lian, C., D. Zhao, and J. Cheng. 2003. "A Fuzzy Logic Based Evolutionary Neural Network for Automotive Residual Value Forecast." In *International Conference on Information Technology: Research and Education, 2003. Proceedings. ITRE2003*, 545–48. IEEE.



- Rashed, A., S. Jawed, J. Rehberg, J. Grabocka, L. Schmidt-Thieme, and A. Hintsches. 2019. "A Deep Multi-task Approach for Residual Value Forecasting." In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, 467–82. Springer.
- Vinutha, H., B. Poornima, and B. Sagar. 2018. "Detection of Outliers Using Interquartile Range Technique from Intrusion Dataset." In *Information and Decision Sciences*, 511–18. Springer.
- Wu, J.-D., C.-C. Hsu, and H.-C. Chen. 2009. "An Expert System of Price Forecasting for used Cars Using Adaptive Neuro-Fuzzy Inference." *Expert Systems with Applications* 36(4): 7809–17.

How to cite this article: Kim, M., J. Choi, J. Kim, W. Kim, Y. Baek, G. Bang, K. Son, Y. Ryou, and K.-E. Kim. 2023. "Trustworthy residual vehicle value prediction for auto finance." *AI Magazine* 44: 394–405. <https://doi.org/10.1002/aaai.12136>

AUTHOR BIOGRAPHIES

Mihye Kim is a manager researcher in the risk modeling team of Hyundai Capital Services. She holds master's degree in Computational Science and Engineering (Applied Mathematics) from Yonsei University. With years of experience in risk management and modeling, she specializes in financial data analysis utilizing machine learning techniques. Her research interests include risk modeling, time series analysis, anomaly detection, and deep learning, with a focus on developing accurate and efficient models for financial applications. She enjoys uncovering complex causal relationships hidden beneath phenomena through data analysis.

Jimyung Choi is a manager researcher in the Risk Modeling team at Hyundai Capital Services. Her interests and expertise lie in data analysis and the development of machine learning models. She has background in Statistics and has acquired extensive knowledge in the financial field during her several years tenure in the card industry. Building upon this experience, she has successfully led risk modeling projects at Hyundai Capital, serving as an advisor for her team. Her strong analytical abilities and attention to detail make her a valuable asset to the team.

Jaehyun Kim is the leader of the Personal Finance Credit Team at Hyundai Capital. He is an expert in risk management in various financial industries such as credit cards, auto finance, and personal loans. In particular, he has a strong expertise in risk modeling.

Currently, he is making significant contributions to the company by effectively utilizing machine learning and AI technologies to establish robust risk strategies.

Wooyoung Kim has been leading the modeling team at Hyundai Capital Services for several years. She established the modeling team and has played a leading role in applying various AI/ML methodologies not only within Hyundai Capital but also in the domestic financial market. In her current role, She is focused on not only developing credit rating models but also exploring various cutting-edge technologies including strategy recommendation systems, time-series forecasting models, and clustering-based data analysis. With her wealth of experience in credit rating modeling, she has helped several leading card companies in Korea to improve their modeling techniques. Prior to joining Hyundai Capital, she worked for several leading card companies, where she spent several years performing credit-scoring modeling tasks.

Yeonung Baek is a highly experienced financial risk management professional, currently serving as the chief risk officer(CRO) and senior vice president in the Risk Division of Hyundai Capital. With a successful career spanning multiple industries, including insurance and commercial loans, he has held key executive positions and gained recognition for her expertise in financial risk management. As the leader of hundreds of risk division employees, he excels at managing teams and achieving organizational success. His leadership has been instrumental in implementing innovative risk management strategies and driving Hyundai Capital to be at the forefront of risk management in the financial industry. His commitment to excellence and his passion for his work make her a respected and valued member of the financial community.

Gisuk Bang is a seasoned risk management professional with expertise in the fields of auto finance, personal loans, and credit cards. With years of executive experience at Hyundai Capital, he has played key roles in various areas of the company, including serving as the chief risk officer (CRO). He has excelled in risk management, strategic planning, and sales management, making him an indispensable member of the leadership team. He is highly regarded by his peers and has received recognition for his contributions across the organization. As the vice president of the Personal Finance Department, he continues to lead with distinction and drive the company forward.

Kwangwoon Son is a senior manager and senior technical engineer in the risk modeling team. With over 12 years of experience in IT, he has primarily been involved in designing and implementing architectures for asset financial systems and managing related projects. He has specialized in business process reengineering and automation, which has allowed him to take a leading role in various machine learning projects within the Risk Modeling team. He is known for his technical expertise and his ability to lead cross-functional teams to deliver successful outcomes.

Yoenman Ryou is a manager of the risk modeling team at Hyundai Capital. With over 10 years of experience in risk modeling, he brings a wealth of knowledge and expertise to his role. Prior to joining Hyundai Capital, he worked for several years in the insurance industry, where he gained valuable experience in risk management. Yoenman holds a Master's degree in Statistics from Yonsei University, and is highly skilled in various modeling techniques.

Kee-Eung Kim is a professor with Kim Jaechul Graduate School of AI, KAIST. He received the BS degree in computer science from the Korea Advanced Institute of Science and Technology, Daejeon, South Korea, in 1995, and the ScM and PhD degrees in computer science from Brown University, Providence, RI, USA, in 1998 and 2001, respectively. From 2001 to 2006, he was a Senior Software Engineer with Samsung SDS, South Korea, and a Senior Research staff member with Samsung advanced institute of technology, South Korea. In 2006, he joined the Faculty of Computer Science Department, KAIST. His research interests include representations and algorithms for sequential decision making problems in artificial intelligence and machine learning, including Markov decision processes, and reinforcement learning.